Cloverb: Clock-Aware, Long-History Human Move Prediction with an Efficient Square-Token Transformer

The Cloverb authors
https://github.com/BarioIsCoding/cloverb

October 2026

Abstract

Maia-3 (Chessformer) is the strongest open square-token model for predicting human chess moves. We reproduce its evaluation protocol exactly and its reported accuracies to within a third of a point, benchmark all nine released checkpoints, and audit the paper and code: we find that its public engine runs without the move history it was trained with (costing 2.2 points of accuracy), that its history encoding and Elo conditioning waste most of the input, that its published models ignore the clock, and that its headline ablation omits the baseline needed to credit its dynamic attention bias. We then present **Cloverb**, a square-token transformer conditioned on exactly what a player has at the board: their own rating, the time control, the time left on their clock, and the last 64 plies encoded as moves. A simple prior, initialising static attention biases with chess movement patterns, makes Cloverb learn 5.5 points faster than the Maia-3 architecture at equal data, and in these 512k-position diagnostics a dynamic bias term in the style of Chessformer's adds nothing. With a C data pipeline, a fused AVX-512 attention kernel and bf16 training on Intel AMX, all experiments, including the final model distilled from Maia-3 through amortised teacher pools, ran on 4-core CPUs without a GPU. On the ALLIE benchmark under Maia's protocol the final 9.2M-parameter Cloverb reaches 55.3% top-1 and 82.4% top-3 move matching (Maia-3-5M: 55.3%, Maia-3-79M: 56.8% on the same positions).

1 Introduction

Predicting the move a human will play, rather than the best move, underpins human-like engines, tutoring and analysis tools. The Maia line of work [1, 2] and ALLIE [3] established the task; Chessformer [4] (released as Maia-3) is the strongest open square-token model, at 57.1% move-matching accuracy on the ALLIE test set with 79M parameters; Allie 2.0 [5], a 5.6B-parameter mixture-of-experts, now reports a 0.31-point gain over it under Maia's protocol on 2026 games.

We started by reproducing Maia-3: we downloaded all nine released checkpoints, rebuilt the exact 884,049-position evaluation protocol, and benchmarked every variant (top-1, top-3, cross-entropy). The numbers reproduce to within a third of a point, but reading the paper next to the released code exposed a series of inaccuracies and design decisions that leave accuracy and efficiency on the table (Section 3). Three matter most for a human-move model: (i) the published models never see the *clock* (only the separately released 23M-ponder checkpoint reads it, and the released code cannot load that checkpoint), and the protocol discards every position played under time pressure; (ii) history is eight board snapshots that alternate orientation every ply and are fused by a single linear map, which may be why "more history does not help"; (iii) skill is encoded by concatenating two 128-d interpolated vectors to every one of the 64 tokens, i.e. two scalars occupying 73% of the input.

Cloverb keeps what works in Chessformer (squares as tokens, a from–to policy head) and re-designs the rest around the information a human player actually has and around efficiency:

- **Inputs a player sees.** The player's own Elo (not the opponent's), the time control (base seconds and increment), the time left on the player's clock, and the last ≤ 64 plies as *moves* scatter-added onto the squares they touched, all in one mover-relative frame; castling rights live on the rook squares and the en-passant square is marked.

Table 1: Released Maia-3 checkpoints on 100,016 positions of the ALLIE test set under Maia's protocol (first 10 plies and time pressure removed). "Paper" is the accuracy reported on the full 884,049 positions. 95% half-width ≈0.4 points with positions clustered by game (0.31 if they were independent); differences between checkpoints, paired on the same positions, have half-widths of 0.15–0.19 points.

Checkpoint	Params	Paper top-1	Top-1	Top-3	CE (nats)
Maia-3 79M	78.9M	57.1	56.82	83.51	1.302
Maia-3 23M-ponder (clock)	23.0M	–	56.53	83.32	1.311
Maia-3 23M	22.9M	56.6	56.32	83.09	1.321
Maia-3 5M	5.2M	55.4	55.27	82.33	1.365
Maia-3 3M (ablation)	3.2M	54.8	54.58	81.77	1.389
5M, absolute PE	4.6M	54.7	54.37	81.62	1.395
5M, relative bias	4.6M	54.6	54.47	81.52	1.396
5M, no history	5.2M	54.0	53.79	81.34	1.410
5M, 32 positions	5.3M	55.4	55.15	82.26	1.369
5M, UCI default (history off)	5.2M	–	53.04	80.80	1.437

- **Conditioning where it acts.** Elo and clock features form one vector that modulates every block through adaptive normalisation (adaLN-Zero [7]) instead of being copied 64 times.
- **Geometry as a prior, not a generator.** Every head of every layer owns a static 64×64 attention bias; most are *initialised* with chess movement patterns (rook lines, diagonals, knight jumps, king steps, pawn captures), the others with a same-file pattern, a distance decay, or zero. This one change makes Cloverb learn 5.5 points faster than the Maia-3 architecture at equal data; in the same diagnostics, a cheap dynamic bias in the style of GAB adds nothing on top of it.
- **Systems.** A C chess core that parses Lichess PGN at >0.3M plies/s and builds training batches by replaying games on the fly (3.6 μs/sample, 6 bytes of storage per ply); a fused AVX-512 attention kernel for 64-token sequences; bf16 training on Intel AMX with `torch.compile` (fp32 on CPUs without AMX); jemalloc. Everything in this paper ran on *4-core CPUs*, without a GPU.

2 Background

Maia [1] trained one CNN per rating band; **Maia-2** [2] unified them with skill-aware attention over CNN channels; **ALLIE** [3] is a decoder-only transformer over move sequences with clock and Elo tokens; **Chessformer / Maia-3** [4] uses an encoder over 64 square tokens with Geometric Attention Bias (GAB), an attention bias generated from a compressed board representation mixing shared templates. **Allie 2.0** [5], a 5.6B-parameter mixture-of-experts trained on 75B tokens with eight GPUs, reports a 0.31-point top-1 gain over Maia-3-79M on Maia's protocol and shows that the clock matters most in time trouble. Transformers trained on engine annotations of board states also play strong chess without search [6]; our target is the human move instead.

3 Reproducing and auditing Maia-3

Protocol. The ALLIE test set [3] contains 18,239 Lichess blitz games from 2022. Dropping each game's first 10 plies and every ply from the first one whose mover has fewer than 30 s selects *exactly* the 884,049 positions reported in the paper. We score a fixed sample of 100,016 positions (2,101 whole games) of that protocol with the official model code (bf16 trunk, fp32 policy head; Table 1). Our sample is harder than the full set for every variant, by 0.13–0.33 points (0.23 on average; the paper reports one decimal). All orderings of the paper's results are reproduced except absolute vs. relative position encoding (54.37 vs. 54.47 here, 54.7 vs. 54.6 in the paper; a paired difference of -0.09 ± 0.19 points, i.e. within noise), and all gaps agree to within 0.2 points.

Findings. The full list with evidence is in the repository (`docs/critique.md`); the most consequential are:

1. *The released engine runs without history by default.* `maia3-uci` fills all history slots with the current position unless `--use-uci-history` is passed; this costs 2.2 points of top-1 accuracy, more than training without history at all.
2. *History is alternating and input-only.* Each past board is mirrored to its own side to move, so consecutive frames are flipped relative to each other, and all frames are fused by one linear map. The paper's conclusion that 32 positions are no better than 8 may be a property of this encoding rather than of chess: the measurement reproduces (HIST32 vs. 5M), but it does not separate the two.
3. *No clock, and time pressure is removed from training and evaluation.* The unreleased-in-code MAIA3-23M-PONDER checkpoint, which does read the clock, is better even on the pressure-free protocol (top-1 $+0.21 \pm 0.15$ points, paired and clustered by game; cross-entropy 1.311 vs. 1.321; Table 1), but it takes the clock as raw seconds and collapses at long time controls (38% top-1 at a base time of 25 minutes or more, against 56% for Maia-3-23M; Section 5.3).
4. *Elo as concatenated soft embeddings* is affine in the rating (the paper concedes it is equivalent to a scalar) yet fills 256 of 352 input channels on all 64 tokens.
5. *Missing state*: castling rights and the 50-move counter are not encoded in the human models, and en passant is only implicit in the history frames (lost when history is off, as in the UCI default), although the paper lists all three as part of the Markov state; its own Leela experiments found that such inputs did not change performance.
6. *Accounting*: "FLOPs" are multiply–accumulates (2× off); the parameter counts of the GAB models do not match the checkpoints (5M: 4.91M stated vs. 5.23M; 3M: 2.98M vs. 3.15M), while the absolute-PE and relative-bias models do (4.58M); an undocumented third ("ponder") head exists and one released checkpoint cannot be loaded by the released code; the value head is d wide, not 128; "all code and data" is open-sourced but no training code or data is.
7. *Missing baseline*: GAB is compared to absolute embeddings and to a translation-invariant 15×15 relative bias, but not to a static, per-square 64×64 bias, so the benefit of GAB being *dynamic* is not isolated (the paper's own Table 3 reports 0.77 cross-position correlation of GAB maps).
8. *Evaluation precision*: on CPUs, bf16 policy logits tie often enough that a strict rank count inflates top-1 by about one point (55.3%→56.4% for Maia-3-5M); we compute all policy heads in fp32.

4 Cloverb

4.1 Inputs

All squares are expressed from the mover's point of view (vertical flip and colour swap when Black moves). Token s starts as

$$x_s = E_{\text{piece}}[p_s] + E_{\text{sq}}[s] + \textstyle\sum_r \mathbf{1}[\text{right } r]\,\mathbf{1}[s = \text{rook}(r)]\,E_{\text{castle}}[r] + \mathbf{1}[s = \text{ep}]\,E_{\text{ep}} + \sum_{j=1}^{64} \big(\mathbf{1}[s = f_j]F_j + \mathbf{1}[s = t_j]T_j\big),$$

where (f_j, t_j) are the from/to squares of the move played j plies ago, $F_j = E_{\text{from}}[j] + E^{\text{from}}_{\text{pc}}[\text{side}_j, \text{piece}_j]$ and $T_j = E_{\text{to}}[j] + E^{\text{to}}_{\text{pc}}[\text{side}_j, \text{piece}_j] + \mathbf{1}[\text{capture}_j]\,E_{\text{cap}}[\text{side}_j, \text{captured}_j]$, with captured_j the type of the piece taken (the diagnostic baseline of Section 5.2 uses $E_{\text{cap}}[\text{side}_j]$ only). The history term is a scatter-add of 128 vectors per position: it costs $O(64d)$ instead of a $(12n{+}256) \times d$ projection per token, it is lossless about the move sequence, and it lets the *same* square token accumulate everything that happened on it.

The conditioning vector $c = \text{MLP}(\phi)$ is built from: a hat-function (linear spline) encoding of the player's Elo on 15 anchors (400–3200), a hat encoding of $\log(1 + \text{time left})$ on 14 anchors (2 s–1 h) together with the same quantity as a scalar, log base, log increment, time left relative to base + 40 inc, the halfmove and fullmove counters, castling rights, check, and the number of available history plies. A linear projection $W_c c$ is added to every token, and c drives adaLN-Zero modulation (scale, shift and residual gate for attention and MLP) in every block.

4.2 Trunk and heads

Each block applies attention with RMS-normalised queries and keys [8] and an additive bias $B_{\ell h} = S_{\ell h} + \sum_{t=1}^{T} a_{\ell h t}(\bar{x}, c)\, U_t$: $S_{\ell h} \in \mathbb{R}^{64\times 64}$ is a static learned bias per layer and head, and the optional second term mixes T templates shared by all layers with coefficients from the mean token and c (a cheap GAB analogue). The static biases are initialised with chess geometry: head 1 gets +2 on rook lines, head 2 on diagonals, head 3 on knight jumps, head 4 on king steps, head 5 on pawn-capture squares, head 6 on the same file, head 7 a $-0.5\cdot$Chebyshev-distance locality term, head 8 starts at zero (Section 5.2 shows that this prior is what matters and that T=0 suffices). The MLP is $W_2\,\mathrm{ReLU}(W_1 h)^2$ [9]. Blocks are post-norm, as in Chessformer: $x \leftarrow \mathrm{RMSNorm}(x + (1+g_1)\,\mathrm{Attn}(x(1+s_1)+b_1))$ and likewise for the MLP, with (s, b, g) from adaLN-Zero; output projections are zero-initialised.

The policy head scores move $(u \to v)$ as $\langle W_f x_u, W_t x_v\rangle/\sqrt{d}$ over Maia-3's 4352-move vocabulary (64×64 plus 256 promotions, whose piece bias comes from the destination token), so Maia-3 checkpoints can serve as teachers. Two auxiliary heads read $[\bar{x}, c]$: game outcome (WDL) and think time (32 log-spaced buckets), the latter giving the engine a human-like time-management signal. The training loss is $\mathcal{L} = \mathrm{CE}_{\mathrm{legal}}(\mathrm{move}) + 0.1\,\mathrm{CE}(\mathrm{WDL}) + 0.05\,\mathrm{CE}(\mathrm{time})$, with the softmax restricted to legal moves. With a teacher [10], the move term becomes $(1-w)\,\mathrm{CE}_{\mathrm{legal}}(\mathrm{move}) + w\,\mathrm{CE}_{\mathrm{legal}}(p_{\mathrm{teacher}})$ (Section 5.3); $w = 0$ in every comparison without a teacher.

4.3 Systems

Data. Games are stored whole: per ply a 16-bit move, the 16-bit clock before the move and the 16-bit time spent; per game ratings, time control and result. A C core parses Lichess SAN movetext with clock comments (its move generation and position encoding are tested against `python-chess` on random games, and on the ALLIE test set it selects exactly the 884,049 positions of Maia's protocol), and builds batches by replaying each sampled game to the sampled ply, emitting the board, castling/en-passant state, the 64-ply history and a legal-move bitmask in ≈3.6 μs per sample on one core, in a background thread (ctypes releases the GIL). 5.5M games (372M plies) occupy 2.2 GB and any position can be sampled in any epoch.

Compute. Training uses bf16 autocast on Intel AMX with `torch.compile`. Profiling showed that GEMMs were only a third of a step: attention over $B \cdot H$ tiny 64×64 problems and memory-bound pointwise kernels dominated. PyTorch's fused CPU attention reached only ≈26 GFLOP/s on this shape, so we wrote a fused AVX-512 kernel that keeps each head's 64×64 scores in L1, fuses the static and template biases, and recomputes scores in the backward pass as FlashAttention does [12] (custom `torch.library` op, exact to $2 \cdot 10^{-6}$ against PyTorch in fp32). Replacing glibc malloc by jemalloc removed page-fault storms on large tensors (+30% with our first kernel; once the projection layout and QK normalisation were fused into the kernel, the two allocators were within 1%). jemalloc must then run with a single arena: with per-thread arenas and decay disabled, retained memory under the compiled step grew by ≈0.3 GB per minute, and the first attempt at the final run was killed for lack of memory after about 30 minutes.

5 Experiments

Data. Training uses four Lichess parquet shards [11] from August and September 2025 (after Maia-3's training window, January 2023–July 2025): 5.54M rated games with clocks from bullet to classical, 372M plies. A position is drawn by picking a (game, side) pair with weight $p(\mathrm{Elo\ bin})^{-0.3} \times w_{\mathrm{tc}}$ (bullet 0.6, blitz 1, rapid 1.3, classical 4) times the number of moves that side played, then a uniform ply of that side; the first 10 plies of each game (moves 1–5) are kept with probability 0.5. Unlike Maia-3 we keep time-pressure positions.

Test sets. (1) The ALLIE 2022 blitz test set under Maia's protocol (100,016-position sample, Table 1); (2) the same 2,101 games with *all* 141,390 plies, i.e. including openings and time pressure; (3) fresh games from September 2026 (120k games streamed from the official dump, never seen by any model): 25k positions each of bullet, blitz, rapid and classical (all plies), and 50k blitz positions under Maia's protocol. During training we monitored a fixed 10k-position subset of (1) (from 2,003 of its 2,101 games). It decided the switch to

Table 2: Equal-data comparison on the ALLIE 10k sample (EMA weights): top-1 (%) / cross-entropy after N training positions, same data stream and evaluation for all runs; throughput is end-to-end training positions per second on a 4-core CPU.

Run	pos/s	0.51M	1.02M	1.54M
Chessformer-5M (Maia-3 architecture)	299	38.8 / 2.09	43.6 / 1.85	45.9 / 1.74
Cloverb d256 L8 (no teacher)	306	37.5 / 2.14	41.7 / 1.94	44.7 / 1.82
Cloverb + online Maia-3-23M teacher	167	42.9 / 1.90	–	–

Table 3: Single-factor diagnostics after 512k training positions (EMA weights, ALLIE 10k sample; same data stream for every row). Indented rows each change one factor of the geometry-initialised model.

Configuration	Top-1	Top-3	CE
Chessformer-5M (authors' module)	38.8	63.4	2.088
Cloverb, zero-initialised biases	37.5	61.5	2.140
Cloverb + geometry initialisation	44.4	71.1	1.823
+ post-norm blocks	45.0	72.4	1.791
+ captured-piece type in history	44.7	71.5	1.818
without dynamic templates (T=0)	44.4	71.1	1.823

pure distillation and the start of the learning-rate decay, so (1) and (2) are not strictly held out for the final model; the fresh 2026 sets (3) are. The choice between the raw and the averaged (EMA) final weights was made on 10k protocol positions from other ALLIE games, in none of the test samples.

5.1 Equal-data comparison with Chessformer

We trained Cloverb (d=256, 8 layers, MLP×2, FLOP-matched to Maia-3-5M: 0.30 vs. 0.32 GMAC per position) and the authors' own Chessformer-5M module (inputs as in `maia3.uci` with `--use-uci-history`: the last 8 boards and both Elos) on the identical stream of 1.54M positions with the same optimiser, schedule shape and evaluation (Table 2). Two findings shaped the final system.

First, Chessformer's recipe is the bottleneck, not its architecture. With a sensible learning rate, preprocessing and kernels, the published 5M architecture learns fast: 45.9% after 1.5M positions, i.e. 81% of the way from chance (a uniform guess among legal moves, 6.1%) to its published accuracy after 0.3% of its published training data. The first Cloverb configuration learned *more slowly per position* (44.7%) even though it trains about 2% faster per second; Section 5.2 isolates why.

Second, a teacher is worth more than data. Distilling [10] the released Maia-3-23M online (soft targets with weight 0.7 plus the human move) lifts Cloverb from 37.5% to 42.9% at 0.5M positions, beyond what the student reaches with twice the data. Because the teacher's forward pass costs about as much as a student training step, we amortise it: the final model trains for several epochs on a fixed pool of positions labelled once by the clock-aware Maia-3-23M-ponder checkpoint (top-16 moves, which hold 99.2% of the probability mass).

5.2 Diagnosing the per-sample gap: geometry is the missing prior

We ran single-factor diagnostics for 512k positions each on the identical data stream (Table 3). One change dominates everything else: *initialising the static attention biases with chess geometry.* Each of the first heads of every layer starts with a +2 logit bonus on the squares reachable by one movement pattern (rook lines, bishop diagonals, knight jumps, king steps, pawn-capture squares, same file) and one head with a Chebyshev-distance penalty; the biases remain fully trainable. This single prior lifts Cloverb from 37.5% to 44.4% after 512k positions, *5.5 points above the Maia-3 architecture* at equal data (44.37% vs. 38.84%; that run used Chessformer's lower peak learning rate, 5×10^{-4}), and nearly matches (44.4% vs. 44.7%) what the zero-initialised model reached with three times the data. It bears on the paper's missing baseline: early in training, much of what a dynamic bias could contribute is geometry, which a static, per-square bias with the

right initialisation supplies at a fraction of the cost (we did not retrain Chessformer itself with such a bias). Post-norm blocks (as in Chessformer) add a further 0.6 points on top of the geometric prior; encoding *which* piece was captured in the move history adds 0.3 points for one extra embedding lookup; and the dynamic template mix contributes nothing (44.35% without vs. 44.37% with, equal cross-entropy), so the final model drops it and keeps only static, geometry-initialised biases.

5.3 Final model

Recipe. The final model has 9.22M parameters: $d = 256$, 10 post-norm blocks with 8 heads, ReLU^2 MLPs of width $3d$, geometry-initialised static attention biases, no dynamic templates, and the captured piece in the move history. Six pools of 2M positions each, drawn with the training sampler, were labelled once by Maia-3-23M-ponder (top-16 moves; about 85 minutes per pool at 390 positions/s). Training uses AdamW ($\beta = (0.9, 0.98)$, weight decay 0.05, learning rate 10^{-3}, 1k warm-up steps), batch 512, gradient clipping at 1, and an EMA of the weights with decay 0.9998. Steps 0–74.5k ran with bf16 autocast on Intel AMX at about 240–270 positions/s; after the job moved to a 4-core Xeon without AMX, the last 5.5k steps (6.9% of training) ran in fp32 without autocast, at about 90 positions/s, with the q/k RMSNorm keeping the bf16 build's ϵ so that the network computes the same function (KL $1.7 \cdot 10^{-5}$ nats between the two precisions).

Switching to pure distillation mid-run. For the first 24k steps the target mixed the teacher's distribution (weight 0.7) with the human move (0.3); below 30 s on the mover's clock, where the teacher never trained, the mix was 0.4/0.6. Projecting the learning curve showed this would land near Maia-3-5M, so at step 24k we branched two 3k-step learning-rate cool-downs from the same checkpoint on identical batches. With the human move removed outside time trouble, the cooled-down raw (non-EMA) weights reached a cross-entropy of 1.3970 and 54.36% top-1, against 1.3996 and 54.19% with the mixed target (EMA weights: 1.4032 and 54.10% vs. 1.4045 and 54.09%). The raw cross-entropy was already 0.008 lower halfway through the cool-down, consistent with less gradient noise. The run continued from step 24k with pure teacher targets (still 0.4/0.6 under 30 s) to 80k steps. By step 48k the constant-rate phase gained only 0.001 nats per 2k steps, so the linear decay to 2% of the peak rate was started there, covering the last 40% of training instead of the planned 30%. In total the model sees about 41M samples, 3.4 passes over the pool.

How far distillation can go. A student trained to match a teacher approaches the teacher's accuracy from below. Table 5 scores the four main Maia-3 checkpoints and ensembles of them on the 10k development subset of the ALLIE test sample. The clock-aware teacher we use is 0.2 points below Maia-3-79M; the best ensemble (79M, 23M and 23M-ponder, averaged logits) is 0.3 points above it. With our single teacher, distillation alone is therefore bounded about 0.2 points below Maia-3-79M under Maia's protocol; only an ensemble teacher could edge past it. Beating it clearly needs what a teacher cannot supply: far more human moves, or information the teacher does not see. Cloverb's clock, increment and 64-ply history are such information, and they matter most in exactly the positions Maia's protocol removes.

A teacher outside its range. The fresh 2026 games exposed a failure that the blitz-only ALLIE set cannot show: on classical games the 80k-step model matched only 40.2% of moves, against 53.7% for Maia-3-5M, and its deficit grew steadily with the base time. To locate the cause without touching the test sets we held out 480 games of the training months, 120 per base-time bucket, and scored both teachers on all their plies (Table 4). The clock-aware Maia-3-23M-ponder collapses at long time controls (38.0% at 25 minutes or more, against 55.6% for Maia-3-23M, which does not read the clock) and is already worse at 10 and 15–20 minutes; it is only marginally better in blitz. Maia-3 is trained on blitz games only [4], and the ponder model's time inputs are raw seconds divided by 100, so a 30-minute game feeds it values about ten times larger than any it saw in training; the models without a clock input cannot be pushed out of range this way. Cloverb, trained with pure teacher targets, had faithfully learned the collapse. We therefore kept the ponder labels only for games with a base time under 10 minutes (78.5% of the pool entries), labelled 2.6M positions of games with a base time of 10 minutes or more with Maia-3-23M (the held-out games excluded), and continued training the final weights on both pools for 8k steps with a fresh schedule (peak learning rate $3 \cdot 10^{-4}$, 300 warm-up steps, linear decay over the second half to 2%, EMA decay 0.999, other settings

Table 4: Top-1 (%) by base time on 480 held-out games of the training months (120 per column, all plies; 8,559, 8,514, 8,392, 8,247 positions), kept out of the fine-tuning pools (the original pools hold 2,280 of their positions among 12M entries). The clock-aware teacher collapses at long time controls; the fix relabels games with a base time of 10 minutes or more with Maia-3-23M.

Model	3–5 min	10 min	15–20 min	25+ min
Maia-3 23M-ponder (teacher)	57.4	54.4	52.3	38.0
Maia-3 23M (teacher)	57.1	55.3	54.8	55.6

Table 5: Maia-3 checkpoints and their ensembles as teachers, on the 10k-position development subset of the ALLIE test sample (Maia protocol). Ensembles average the members' log-probabilities ("logits") or probabilities.

Model	Top-1	Top-3	CE
Maia-3 79M + 23M + 23M-ponder (logits)	56.94	84.22	1.288
Maia-3 79M + 23M-ponder (logits)	56.84	84.16	1.285
Maia-3 79M	56.66	84.36	1.290
Maia-3 23M	56.58	83.74	1.307
Maia-3 23M-ponder (our teacher)	56.49	83.86	1.299
Maia-3 5M	55.25	82.75	1.351

unchanged). Raw or EMA weights were again chosen by cross-entropy on held-out positions only (the 10k ALLIE selection positions plus the 480 held-out games).

5.4 Inference

Cloverb ships as a standalone C program (one translation unit, two source files) with no Python or PyTorch dependency at run time: the C chess core encodes the position and its history, the trunk's GEMMs run on AMX bf16 tiles with weights packed once into VNNI tiles, attention runs in AVX-512 with each head's scores kept in registers and L1, and the policy head stays in fp32. On 500 random positions of the fresh 2026 games, an fp32 build (AVX-512, no AMX) reproduces the PyTorch forward pass of the trained weights to within $2.2 \cdot 10^{-6}$ in every legal-move probability and the AMX bf16 build to within $2.6 \cdot 10^{-3}$ (mean KL $2.4 \cdot 10^{-6}$ nats), both with the same top move on every position (`results/engine_check.json`).

6 Limitations

- **Compute.** Every number here comes from 4-core Xeon CPUs, one at a time, over a few days (the container moved between hosts: hosts with AMX ran most of training, the Maia-3 benchmarks and the speed measurements; a host without AMX ran the last 5.5k training steps and the final Cloverb evaluation, both in fp32). Cloverb saw tens of millions of positions, Maia-3 half a billion. Our conclusions about architecture are about learning speed in this regime; at Maia-3's budget the geometric prior's advantage may shrink.
- **Teacher.** The final model is distilled from a released Maia-3 checkpoint (with the human move as a second target), so its accuracy partly inherits Maia-3's; the equal-data comparisons in Table 3 and the first two rows of Table 2 use no teacher. Maia-3 never trained on positions with less than 30 s on the clock, so in time trouble the human move keeps a weight of 0.6 in our loss (elsewhere 0.3 for the first 24k steps and 0 afterwards).
- **Opponent rating.** By design Cloverb sees only the mover's rating, as an engine user would set it; Maia-3 also sees the opponent's.
- **Test sets.** The ALLIE test set is blitz from 2022; our training games are from 2025. The fresh 2026 test set covers all time controls, but only 615 of its 120k games are classical (390 of them are scored), and each all-plies set comes from fewer than 400 games.

Table 6: ALLIE 2022 blitz test set. Left: Maia's protocol (100,016 positions; first 10 plies and time pressure removed). Right: every ply of the same 2,101 games (141,390 positions), including openings and time trouble. Maia-3 sees both ratings; Cloverb only the mover's.

	Maia protocol			All plies		
Model	Top-1	Top-3	CE	Top-1	Top-3	CE
Maia-3 79M	56.82	83.5	1.302	–	–	–
Maia-3 23M-ponder	56.53	83.3	1.311	–	–	–
Maia-3 23M	56.32	83.1	1.321	–	–	–
Maia-3 5M	55.27	82.3	1.365	54.93	82.3	1.372
Cloverb (9.2M)	55.28	82.4	1.363	55.14	82.5	1.365

Table 7: Fresh games from September 2026 (unseen by every model), top-1 / top-3 (%). All-plies sets are every ply of whole games, about 25k positions per set from 358–399 games each; the Maia-protocol set has 50k positions from 958 games. Positions of one game are correlated, so differences below about one point are within noise.

Model	Blitz (Maia prot.)	Bullet	Blitz	Rapid	Classical
Maia-3 79M	– / –	– / –	– / –	– / –	– / –
Maia-3 23M-ponder	– / –	– / –	– / –	– / –	– / –
Maia-3 23M	– / –	– / –	– / –	– / –	– / –
Maia-3 5M	– / –	– / –	– / –	– / –	– / –
Cloverb (9.2M)	55.78 / 83.2	53.33 / 80.0	56.01 / 83.3	54.02 / 81.5	40.20 / 66.7

7 Conclusion

Reproducing Maia-3 confirmed its accuracy and exposed avoidable losses in its inputs, defaults and evaluation. Cloverb reorganises the same square-token idea around what a human sees at the board (own rating, clock, the moves that led here) and around the hardware it runs on, and shows that a simple geometric prior on attention, not a dynamic bias generator, is what lets such a model learn chess quickly.

References

[1] R. McIlroy-Young, S. Sen, J. Kleinberg, A. Anderson. Aligning superhuman AI with human behavior: chess as a model system. KDD 2020.

[2] Z. Tang, D. Jiao, R. McIlroy-Young, J. Kleinberg, S. Sen, A. Anderson. Maia-2: a unified model for human-AI alignment in chess. NeurIPS 2024.

[3] Y. Zhang, A. Jacob, V. Lai, D. Fried, D. Ippolito. Human-aligned chess with a bit of search. ICLR 2025.

[4] D. Monroe, G. Eilender, P. Chalmers, Z. Tang, A. Anderson. Chessformer: a unified architecture for chess modeling. ICLR 2026. arXiv:2605.19091.

[5] Y. Zhang et al. Allie 2.0. https://github.com/y0mingzhang/allie, 2026.

[6] A. Ruoss et al. Amortized planning with large-scale transformers: a case study on chess. NeurIPS 2024.

[7] W. Peebles, S. Xie. Scalable diffusion models with transformers. ICCV 2023.

[8] M. Dehghani et al. Scaling vision transformers to 22 billion parameters. ICML 2023.

[9] D. So et al. Primer: searching for efficient transformers for language modeling. NeurIPS 2021.

[10] G. Hinton, O. Vinyals, J. Dean. Distilling the knowledge in a neural network. arXiv:1503.02531, 2015.

[11] Lichess open database. https://database.lichess.org.

Table 8: What the final model uses, measured by removing inputs at evaluation time on the ALLIE sets (the clock is replaced by another position's fraction of base time left).

Evaluation	Top-1	Top-3	CE
Cloverb (full inputs)	55.28	82.4	1.363
history cut to 8 plies	54.81	82.0	1.384
no history	50.38	78.6	1.557
clock shuffled	55.04	82.2	1.371
Cloverb, all plies	55.14	82.5	1.365
clock shuffled	54.81	82.1	1.386

Table 9: Think time and outcome, all plies with a known time spent. Mean absolute error of $\log(1+\text{seconds})$ for Cloverb's median predicted think time, for the best constant (the test set's own median, an oracle) and for the rule of thumb "1/40 of the clock plus the increment"; correlation of predicted and actual log time; share of moves predicted within a factor of two. Right: accuracy and cross-entropy of the win/draw/loss head against the game result.

	Log-time MAE					Outcome	
Test set	Cloverb	Constant	Rule	Corr.	$\leq 2\times$ (%)	Acc. (%)	CE
ALLIE 2022 blitz	0.438	0.630	0.696	0.68	70.9	61.5	0.764
Fresh bullet	0.337	0.458	0.523	0.64	72.6	62.0	0.696
Fresh blitz	0.430	0.611	0.713	0.67	71.4	62.1	0.747
Fresh rapid	0.529	0.733	0.981	0.65	65.7	64.5	0.695
Fresh classical	0.679	0.892	1.487	0.61	55.7	65.8	0.730

[12] T. Dao et al. FlashAttention: fast and memory-efficient exact attention with IO-awareness. NeurIPS 2022.

Table 10: Top-1 (%) by game phase on all plies: openings (first 10 plies), normal play, and time trouble (later plies with under 30 s on the mover's clock, which Maia's protocol and training data exclude). ALLIE: 2022 blitz (21,010, 106,609, 13,771 positions); fresh: the four September 2026 all-plies sets pooled (15,304, 74,061, 10,742).

	ALLIE, all plies			Fresh 2026, all plies		
Model	Open.	Normal	<30 s	Open.	Normal	<30 s
Maia-3 5M	50.83	55.44	57.20	–	–	–
Cloverb (9.2M)	50.56	55.51	59.25	45.28	51.37	55.55

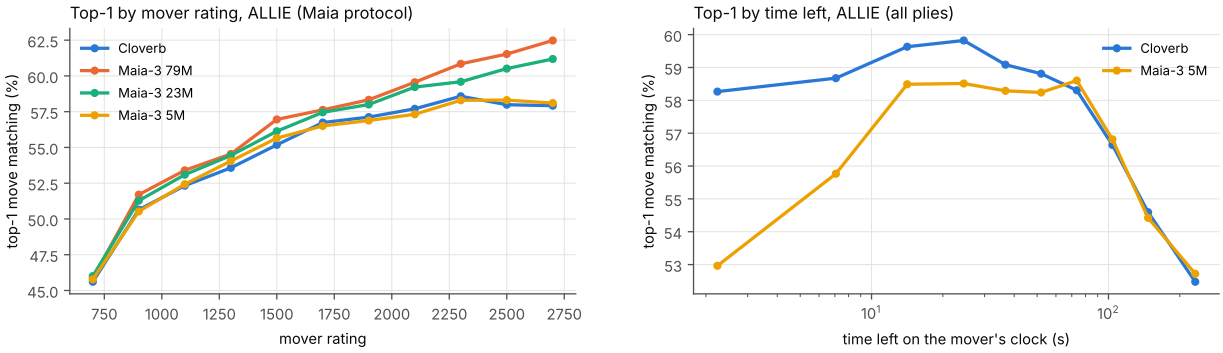


Figure 1: Left: top-1 by mover rating on the ALLIE set (Maia protocol). Right: top-1 by time left on the mover's clock on all plies of the same games; Maia-3 never trained below 30 s.